


Value Kaleidoscope: Engaging AI with Pluralistic Human Values, Rights, and Duties

Taylor Sorensen, Liwei Jiang, Jena D. Hwang , Sydney Levine, Valentina Pyatkin, Peter West, Nouha Dziri, Ximing Lu, Kavel Rao, Chandra Bhagavatula, Maarten Sap, John Tasioulas, Yejin Choi

Motivation

- People weigh multiple, sometimes conflicting, values in everyday decisions (e.g., honesty vs. kindness).
- Standard ML tends to average away disagreement, obscuring meaningful value conflicts.
- **Goal:** Make implicit value tradeoffs explicit so AI can better model human diversity in judgment.
- Represent different normative values instead of a super value.



Contribution

- ValuePrism: 31k situations paired with 218k contextualized values/rights/duties, distilled from GPT-4 and human-validated.
- KALEIDO: An open, multi-task model that generates values, predicts relevance and valence, and explains links to context.
- Systems: KALEIDO SYS (value-set generation/selection) and KALEIDO DEC (interpretable decision aggregation).
- Evidence: Human preferences of Kaleido over GPT-4 outputs; transfer to other frameworks; zero-shot decision reasoning.

ValuePrism: Dataset

- 31k human-written situations; 218k contextualized value/rights/duty items.
- Distilled via GPT-4, then evaluated by crowd workers; 91% items rated high-quality.
- Diverse annotator study to probe “whose values” are represented.
- Released resources: dataset, models, code, and demo.

Type	Total	Unique	Avg. Per Situation
Situations	31.0k	31.0k	1
Values	97.7k	4.2k	3.15
Rights	49.0k	4.6k	1.58
Duties	71.6k	12.8k	2.31

Table 2: VALUEPRISM Dataset Statistics. The total, unique and average per situation statistics of generated values, rights, and duties are shown.

KALEIDO: Tasks and Outputs

- Generation: Propose relevant values/rights/duties for a situation.
- Relevance (binary): Is the item pertinent to this situation?
- Valence (3-way): Does it support, oppose, or depend?
- Explanation: Natural-language rationale linking item to the situation.
- Outputs include textual rationales and scalar scores for relevance/valence.

Situation: Telling a lie to protect a friend's feelings		
Task	Input	Output
Generation	{situation}	Value: Honesty
Generation	{s}	Value: Friend's well-being
Relevance	{s}, Value: Honesty	Yes
Relevance	{s}, Value: Economic well-being	No
Valence	{s}, Value: Honesty	Opposes
Valence	{s}, Value: Friend's well-being	Supports
Explanation	{s}, Value: Honesty	If you value honesty, it may be better to tell the truth even if it hurts feelings.

Table 1: Illustrative examples of each task, with {situation}/{s} standing in for the example situation.

System 1: KALEIDO SYS (Curation Pipeline)

- Step 1: Over-generate candidate values/rights/duties for a situation.
- Step 2: Filter by model-predicted relevance; assign relevance and valence.
- Step 3: Deduplicate similar items; return diverse, high-coverage value sets.
- Tunable precision–recall: generate concise or broad sets depending on need.




Figure 2: KALEIDO^{SYS} system workflow that includes 1) generating 100 values, rights and duties; 2) filtering by relevance as rated by KALEIDO; 3) removing repetitive items; and computing relevance and valence scores for each value, right, and duty.

System 2: KALEIDO DEC (Interpretable Aggregation)

- Aggregate value considerations into an action judgment by weighted linear combination of item valences: $\text{Score}(\text{action}) = \text{sum over items of Relevance} \times \text{Valence} \times (\text{optional user weight})$.
- Interpretable: users can inspect and reweight values; steerable by zeroing or boosting select values.
- Entropy of outputs serves as a proxy for ambiguity/controversialness.

Human Evaluations (Generation Quality)

- Against GPT-4, human raters prefer KALEIDO SYS outputs on correctness and coverage; larger models improve further.
- Annotators struggled to find missing considerations (>99% recall in spot checks), suggesting broad coverage in practice.
- Average items per situation adjustable (e.g., around 8–11 in studies).

Model	Overall	Cover.	Acc.	Avg. #
KALEIDO ^{SYS} 3B	55.5	65.1	58.9	8.2
-relevance	51.9	81.4	64.3	11.2
-text similarity	50.0	60.5	52.9	8.2
verbose	58.0	86.1	69.0	11.1
concise	39.0	27.4	32.4	5.0
KAL ^{SYS} 11B	58.3	71.1	62.5	8.3
KAL ^{SYS} 770M	57.9	67.3	60.8	8.2
KAL ^{SYS} 220M	44.9	59.0	50.8	8.1
KAL ^{SYS} 60M	32.0	53.0	37.1	8.5
Direct Output	42.5	37.9	40.0	6.8
GPT-4	50.0	50.0	50.0	7.0
GPT-3.5-turbo	39.5	49.0	39.8	8.0

Table 5: The overall, coverage and accuracy win rate percentage against GPT-4 by human evaluators along with the average number of generated values, rights, and duties. (Here and throughout, best results within 1% are bolded.)

Human Evaluations (Explanation, Valence, Relevance)

- Explanations and valence judgments from KALEIDO approach GPT-4 quality; best models sometimes slightly exceed GPT-4 on explanations.
- Relevance scores (trained on synthetic labels) correlate positively with human judgments ($r \approx 0.25\text{--}0.31$), indicating useful signal despite proxy training.

Model	Explanation	Valence	Rel. corr.
KALEIDO 3B	92.6	92.0	0.30
KAL 11B	94.8	92.6	0.25
KAL 770M	90.3	90.3	0.31
KAL 220M	86.9	86.3	0.30
KAL 60M	75.9	72.3	0.28
GPT-4	94.7	93.1	-

Table 6: Human Evaluation. Explanation and Valence scores are correctness rates of the output, while Relevance is the correlation of relevance score with the percentage of people who marked a value as relevant.

Generalization to Other Ethical Frameworks

- Zero/few-shot tests on ETHICS subsets (Justice, Deontology, Virtue, Utilitarianism, Commonsense).
- KALEIDO (3B) matches or exceeds ChatGPT on several subsets despite being smaller, especially in Utilitarian and Virtue tasks.
- Suggests learned value representations transfer beyond the trained schema.

Subset	KALEIDO	ChatGPT	Random
Justice	17.5 / 13.3	17.6 / 13.4	6.3 / 6.3
Deont.	19.8 / 15.1	20.6 / 13.8	6.3 / 6.3
Virtue	33.1 / 22.2	24.9 / 22.0	8.2 / 8.2
Util.	76.5 / 66.6	59.4 / 55.1	50.0 / 50.0
Comm.	71.5 / 64.7	80.3 / 68.8	50.0 / 50.0
Average	43.7 / 36.4	40.6 / 34.6	24.2 / 24.2

Table 7: ETHICS few-shot performance. First/second number of each entry is performance on the test/hard test sets respectively. KALEIDO is zero-shot, ChatGPT is few-shot.

Zero-shot Decision Performance on Commonsense Norms

- Applied KALEIDO DEC to moral acceptability tasks (e.g., SBIC, ETHICS, Moral Stories, SocialChem).
- Zero-shot performance exceeds majority class baselines; simple calibration further improves results.
- Still below supervised systems (e.g., Delphi), but interpretable and data-efficient.

Model	SBIC	ETH.	MoSt	SoCh
KALEIDO ^{DEC}	64.4	77.9	75.4	48.2
+label calibration (improvement)	69.3 (+4.9)	78.0 (+0.1)	76.2 (+0.8)	63.0 (+14.8)
Majority class	63.1	51.6	50.0	46.7
Random	33.3	50.0	50.0	33.3
Delphi (SFT)	82.9	86.2	86.5	78.0

Table 8: Zero-shot Performance on COMMONSENSE NORM-BANK: Moral Acceptability.

Entropy Signals Ambiguity

- Higher KALEIDO DEC entropy predicts high-ambiguity or high-controversialness scenarios in external datasets.
- Practical cue: flag cases where judgments likely vary by value priority.




Figure 3: The output entropy of KALEIDO^{DEC} is predictive of ambiguity in MoralChoice and controversialness in SocialChem. A threshold is chosen to maximize F1-score.

Sensitivity to Context

- Small changes in scenario details shift relevance/valence assignments.
- Illustrates contextual grounding: values are not static labels; they respond to situational nuances.

Limitations and Cautions

- Distillation inherits biases of teacher models and prompts; representativeness depends on data and annotator pool.
- Relevance training uses synthetic labels; calibration to human judgments can be improved.
- Interpretable aggregation is linear; complex moral reasoning may require richer composition or deliberation.

Mind the Value-Action Gap: Do LLMs Act in Alignment with Their Values?


Hua Shen, Nicholas Clark, Tanu Mitra

Motivation

- LLMs' stated values may not predict their actions. The discrepancy is the "Value-Action Gap," well-studied in social psychology but under-examined in LLMs.
- Real-world risks: misalignment can amplify stereotypes, bias decisions, or contradict safety expectations.
- Research question: To what extent do LLM-generated value statements align with their value-informed actions?

Framework overview — ValueActionLens

- Contextualization across 12 countries × 11 social topics = 132 scenarios.
- Two tasks: Task 1 elicits value inclinations; Task 2 asks models to choose between value-consistent actions.
- Alignment metrics: alignment rate (F1), alignment distance (L1), alignment ranking.



Dataset: Value-Informed Actions (VIA)

- 14,784 value-informed actions with explanations
- 12 countries, 11 social topics, 56 values

Features	Count	Details or Examples
Countries	12	United States, India, Pakistan, Nigeria, Philippines, United Kingdom, Germany, Uganda, Canada, Egypt, France, Australia
Social Topics	11	Politics, Social Networks, Inequality, Family, Work, Religion, Environment, National Identity, Citizenship, Leisure, Health
Values	56	Social Power, Equality, Choosing Own Goals, Creativity, Honest, etc. See a full list of 56 values and definitions in Table 6.
Inclinations	2	Agree, Disagree
Value-Informed Actions with Explanations	14,784 in total	Value-Informed Actions: I make decisions for my family about which health care provider to visit and ensure everyone follows my recommendations . (highlights are explained actions.) Explanations: This action reflects that I possess the value of Social Power because it demonstrates control and dominance over others by taking charge of critical health care decisions and ensuring compliance from my family members.


Example of a value-action gap

Illustrative case (Health in Nigeria; value: Social Power):

Task 1: Model states “Disagree” with Social Power (control over others, dominance)

Task 2: Selects an action indicating control (e.g., deciding for family and ensuring compliance)

Outcome: misalignment between stated value and action



Data generation pipeline and quality control


Human-in-the-loop process includes three steps:

1. Building eight prompt variants (paraphrase, reorder, requirement changes)
2. Expert prompt selection with inter-rater reliability (Cohen's $\kappa = 0.7073$)
3. Cross-cultural human evaluation (Prolific; 27 annotators; majority vote)

Evaluation metrics used are:

Correctness, Harmlessness, Sufficiency (attributions), Plausibility (explanations)

Objects	Value-Informed Actions		Attributions	Explanations
Metrics	Correct	Harmless	Sufficient	Plausible
AI Researchers	0.93125	0.95625	0.9438	1.00
Annotators	0.8778	0.800	0.8889	0.9222



Tasks and prompt designs


- Task 1: State value inclination
 - SVS- and PVQ-style prompts; Likert scale; eight prompt variants averaged per scenario
- Task 2: Select value-informed action
 - Two actions (agree vs. disagree) with order shuffling; eight prompt variants averaged
- Robustness:
 - multiple prompt variants to reduce bias

Alignment measures

- Alignment rate (F1): compares binarized agreement/disagreement from Task 1 and Task 2
- Alignment distance: element-wise Manhattan (L1) distance between Task 1 and Task 2 matrices across values per scenario
- Alignment ranking: ranks values by distance per scenario to surface largest gaps
- Equations:
 - Distance per value-scenario: $D_{ik} = |v_{ik} - a_{ik}|$
 - Averaged by country/topic: $D_{Ck} = (1/|C|) \sum_{i \in C} |v_{ik} - a_{ik}|$


Big-picture results: countries and topics

- Alignment varies notably across models, cultures, and topics.
- Some models align better overall, but none eliminate gaps.
- Gaps are larger for certain values like Independent and Choosing Own Goals, and differ by country/topic.



Cultural contrast via rankings

- Insight: The “largest-gap” values differ by culture.
- Example: Philippines vs United States—“Choosing Own Goals” ranks among largest gaps in the Philippines but not in the U.S.



Cultural Contrast




Figure 5: Comparing the Alignment Ranking of 56 values in Philippines (top) and United States (bottom).

What gaps imply: risk taxonomy

Risk categories from misaligned samples:

- Individual: discrimination, autonomy violation, privacy, psychological harm
- Interaction: misleading explanations, overconfidence, manipulation
- Societal: misinformation, polarization, undermining institutions

Category Level	Risk Type	Count
Individual	Discrimination	334
	Autonomy Violation	42
	Privacy Invasion	4
	Psychological Harm	3
Interaction	Misleading Explanations	1
	Overconfidence	4
	User Manipulation	1
Societal	Misinformation	14
	Polarization	75
	Undermining Institutions	2

Table 5: The value-action risk taxonomy and statistics in the six LLMs’ generations, indicating potential risks in real-world LLM behaviors.

Takeaways and Mitigations

- Evaluate actions, not just statements: include action-choice or free-form behavior tasks in alignment checks.
- Make evaluations scenario-aware and pluralistic: test across cultures, topics, and a broad set of values (not only fairness/harmlessness).
- Prioritize by rankings: fix the largest-gap values per scenario via targeted red-teaming or data curation.
- Leverage explanations: request reasoned rationales; they improve predictability of values from actions and can guide steering.

Limitations and conclusion

Limitations: fixed value ontology (Schwartz), binary action choices, single-turn/static evaluation, specific to English

Conclusion: ValueActionLens offers a reproducible way to quantify value-action gaps across contexts; findings show substantial gaps and motivate scenario-sensitive, behaviorally grounded alignment assessments. Speaker notes: The method is deliberately simple to enable commensurable measures. Future extensions could add free-form action generation, graded intensities, and dialog.

Future of Work with AI Agents: Auditing Automation and Augmentation Potential across the U.S. Workforce

Yijia Shao, Humishka Zope, Yucheng Jiang, Jiaxin Pei, David Nguyen,
Erik Brynjolfsson, Diyi Yang


Motivation

- The paper introduces a worker-centered auditing framework, a Human Agency Scale (HAS) to quantify ideal human involvement, and a new dataset, WORKBank, with 1,500 workers across 104 occupations rating 844 tasks, plus assessments from 52 AI experts.
- I'll talk about the framework, core findings on automation desire, the desire and capability comparison landscape, human agency patterns across occupations, how skills may shift, and what's mismatched between current investment/research and where workers want help.

Motivation

- AI agents are here; impacts uneven across tasks
- Prior work: capability-centric, narrow domains
- This study: worker-centered, task-level, forward-looking

Auditing framework at a glance



Human Agency Scale (HAS)

- H1: AI does it all
- H2: Minimal human input
- H3: Equal partnership
- H4: AI requires human input
- H5: Continuous human involvement essential

	HAS H1	HAS H2	HAS H3	HAS H4	HAS H5
Team Dynamics	AI Agent Drives Task Completion The AI agent takes primary responsibility for task execution with no or minimal human oversight.		Equal Partnership The human and the AI agent collaborate closely throughout the task.	Human Drives Task Completion The human takes primary responsibility for task execution with varying levels of AI assistance.	
Required Human Involvement	AI agent handles the task entirely on its own without your involvement.	AI agent needs your input at a few key points to achieve better task performance.	AI agent and you work together to outperform either alone.	AI agent needs your input to successfully complete the task.	Task completion fully relies on your involvement.
AI Role	Automation AI replaces human capabilities		Augmentation AI enhances human capabilities		
Example Tasks	• Transcribe data to worksheets and enter data into computer. • Run monthly network reports.	• Devise trading, option, or hedge strategies. • Accept payment on accounts.	• Create core game features, including storylines, role-play mechanics, etc. • Compile and analyze experimental data and adjust experimental designs as necessary.	• Coordinate and direct the financial planning, budgeting, procurement, or investment activities. • Design, plan, organize, or direct orientation and training programs.	• Participate in online forums or conferences to stay abreast of online retailing trends, techniques, or security threats.

Data: WORKBank coverage

- 1,500 workers | 104 occupations
844 tasks
- 52 AI experts (capability + feasible HAS)
- Representative sector coverage





Figure 3: **Sector-level distribution of workers in the WORKBank database compared to U.S. workforce statistics from the Bureau of Labor Statistics.** **a**, Comparison between WORKBank worker distribution and the U.S. workforce employment statistics across all sectors (sectors not included in WORKBank marked with an asterisk). **b**, Comparison between WORKBank worker distribution and the U.S. workforce employment statistics limited to the 104 occupations included in our database.

Usage isn't the same as desire

- High-desire occupations underrepresented in usage logs


**C Percentage of Usage on Claude.ai
(Dec 2024–Jan 2025)**




Desire–Capability landscape: 4 zones

- Green Light: high desire, high capability
- Red Light: low desire, high capability
- R&D Opportunity: high desire, low capability
- Low Priority: low desire, low capability


a Automation Desire-Capability Landscape




Computer and Mathematical




Management



Business and Financial Operations




Arts, Designs, and Media



Investment and research are misaligned

- YC investment: 41% mapped to Low Priority + Red Light
- Papers: concentrated in R&D zone but narrow domains

**b Average Number
of Y Combinator
Companies per Task
by Desire–
Capability Zone
(Cut-off Date:
April 28, 2025)**



HAS patterns across occupations

- Workers prefer more agency than experts deem necessary
- H3 (equal partnership) is the most common worker preference

a Distribution of Tasks by Worker-Desired and Expert-Assessed Feasible HAS Levels

		AI Expert-rated Feasible HAS Level				
		H1	H2	H3	H4	H5
Worker-desired HAS Level	H1	32	34	30	7	0
	H2	112	102	63	36	4
	H3	61	114	71	34	7
	H4	19	41	30	22	1
	H5	1	5	11	7	0

What defines high-agency tasks and shifting skills

- H5 tasks: interpersonal communication and domain expertise
- Emerging shift: info-processing ↓; interpersonal/organizational ↑



Figure 7: Comparing skill rankings by average wage and required human agency. Each line

Implications, limitations, takeaways

- Build for H3 collaboration; align with Green Light + R&D zones
- Reskill toward interpersonal/organizational skills

Limitations

- Task set tied to O*NET; U.S.-focused snapshot
- Worker awareness and preference dynamics may evolve

Artificial Intelligence and the Labor Market: A Scenario-Based Approach

Remarks by Michael S. Barr

Introduction

- How will Artificial Intelligence, especially Generative AI (GenAI), affect workers and the labor market?
- Acknowledging the historical debate: Machines replacing workers is not a new concern.
- However, GenAI's capabilities, generating new content, undertaking complex tasks, and accelerating R&D, suggest a potentially more profound impact.

AI as a General Purpose Technology

- AI exhibits characteristics of a "general purpose technology" (GPT) like electricity or computers.
- Characteristics of a GPT:
 - Widespread adoption
 - Complementary progress in downstream applications
 - Ongoing improvement in core technology
 - Historical GPTs have led to dramatic productivity improvements.

AI's Application in Occupations

- AI's impact at the occupational level:
 - **Substitution:** AI replaces human labor in some tasks.
 - **Complement:** AI enhances human labor in other tasks.
 - **Creation:** AI spurs the development of entirely new tasks and occupations.
- Net effects on employment and wages depend on the balance of these forces.
- Pessimistic view: Mass unemployment (e.g., Luddite fears).
- Economists' skepticism: "Lump of labor fallacy" – work is not fixed.

Scenario 1: Incremental Progress

- **Pace:** Brisk but not breathless adoption; initial quick advancements may plateau due to resource constraints (computing, data, energy).
- **Impact:** AI primarily automates some tasks within occupations, augmenting human capabilities.
- **Examples:** AI copilots for programmers, chatbots for customer support, GenAI for legal research, AI safety features in vehicles.
- **Job Outlook:** Most occupations persist in modified, more productive forms. New jobs emerge (e.g., AI managers, human-AI collaboration specialists, ethicists).
- **Economic Outcome:** Jobs remain plentiful, real wages buoyed by productivity, high labor force participation.
- **Challenges:** Painful dislocation and transition for some workers, potential for skill-biased technological change leading to wage inequality. Retraining is crucial.

Scenario 2: Transformation

- **Pace:** Rapid breakthroughs leading to a complete transformation of the economy.
- **Impact:** AI masters nearly all forms of cognitive work, outperforming human counterparts in many activities, including some traditionally human-centric roles (e.g., medical diagnosis, customer support). Advanced robotics substitute for manual and production jobs.
- **Job Outlook:** Significantly smaller role for human labor. New occupations primarily focused on managing the AI-dominated economy. Some jobs persist due to:
 - Difficulty in automation (physical dexterity), Consumer preference for human interaction, Legal and regulatory protections.
- **Economic Outcome:** Vastly expanded societal resources, revolutionary advances in health, increased leisure.
- **Challenges:** Profound social changes, potential for significant decline in employment, widespread skill obsolescence, structural unemployment, and increased income inequality (capital owners benefiting most).

Monitoring the Future

Key Indicators to Monitor:

1. **AI Adoption Rates:** How many businesses are using AI and how it's affecting operations.
2. **Evolving AI Capabilities:** Benchmark tests of AI models against human performance.
3. **Job Market Data:** Analyze job openings and required skills (a leading indicator).
4. **Occupational & Industry Job Growth:** Observe long-term changes in employment statistics.

Challenges: Disentangling AI's effects from other economic factors.

Implications for Monetary Policy

- AI's structural economic changes could affect the Federal Reserve's dual mandate (maximum employment, stable prices).
- Two Prominent Possibilities:
 - **Natural Rate of Unemployment (u^*):** AI could influence u^* through skill mismatch, lengthy unemployment spells, and demographic shifts in the labor force.
 - **Neutral Interest Rate (r^*):** A permanently higher productivity growth rate (under either AI scenario) tends to raise r^* , requiring a higher real interest rate for a given monetary policy stance.
- Both u^* and r^* are unobservable and difficult to estimate in real-time, requiring careful monitoring of economic activity and financial markets.

Conclusion

- AI is poised to profoundly transform the economy.
- Uncertainty remains regarding the speed and extent of this transformation.
- AI will likely boost productivity, accelerate scientific discovery, and fundamentally alter the nature of work.
- These developments have significant implications for society and central bankers.
- Policymakers must decide how to regulate, adapt education/training, and update labor market regulations to ensure broadly shared economic gains.